

A Post-Quantum Relation-Bound Blind Certification Core with Signature-Gated Threshold Decryption

Research construction draft, version 1.4

23 August 2026

Abstract

This document specifies and proves an abstract post-quantum cryptographic core for privacy-preserving, accountable credentials. During an offline issuance phase, a holder obtains a blind signature only after proving that the hidden ticket, the blind-signing request, the holder secret, and a trace ciphertext are mutually consistent. During opening, threshold decryption shares are released only for a validly signed ticket and an independently authorized case. This *signature-gated threshold decryption* rule replaces a public validity proof carried by every online ticket.

The main result is conditional and modular: assuming simulation-extractable zero knowledge, Blind-UOV one-more unforgeability and blindness, and a compact threshold encryption scheme with post-quantum privacy and unique openings, the construction provides issuer unlinkability, certified ciphertext validity, decoder safety, trace correctness, non-frameability, and a gated chosen-ciphertext privacy notion against fewer than the threshold number of opening authorities. The result deliberately does not claim ordinary IND-CCA security for the bare trace encryption. Version 1.2 also instantiates the trace relation as a Niederreiter syndrome plus a SHAKE256/KMAC256 one-time hash DEM, proves information-theoretic uniqueness of the decoded error, and removes a redundant ticket commitment. Version 1.3 freezes a paper-compatible Blind-UOV-Is/Shorter costing profile, fixes a 368-byte payload, gives a reproducible five-block circuit count, and separates the large issuance NIZK from the CAP proof carried by the final Blind-UOV signature. Version 1.4 adds an executable witness builder, a streaming characteristic-two circuit generator, independent SHAKE256/KMAC256 checks, and full-circuit rejection tests. Native Blind-UOV circuit integration, backend simulation extractability, and the threshold implementation remain explicit obligations.

1 Status, scope, and non-goals

This is the standalone cryptographic core. It intentionally excludes satellite AKE, radio transport, FLEO relay behavior, handover, revocation distribution, and concrete parameter optimization. Issuance and opening may use large proofs, large public keys, and large private keys. The constrained online path carries only a ticket payload and a blind signature.

The construction is designed for the following division of trust.

- The issuing HNCC is *honest but curious*: it authenticates the real identity, enforces a common issuance policy, verifies Π_{issue} , and signs only after acceptance, but later tries to link an observed ticket to an issuance session.
- Holders and network attackers may be malicious and quantum polynomial time.
- Fewer than τ of n opening authorities may be corrupted. Any authorized set of at least τ authorities can open a ticket.
- Issuer signing keys and opening decryption keys are independent even when the organizations operating them overlap.

- Per-user visible metadata is forbidden. Version, epoch, policy, expiration bucket, and key identifiers are common to a sufficiently large anonymity set; otherwise they become an issuer watermark outside the cryptographic game.

Claim boundary. The theorems below are complete reductions for the abstract interfaces and games defined here. They are not a claim that the present Blind-UOV or threshold Niederreiter papers already implement every required interface. The Blind-UOV request equations are fixed in Section 3.2; the exact added circuit is counted in Section 8.3 and executed in Section 8.4. The resulting proof-size and prover runtime figures remain paper-anchored estimates rather than implementation benchmarks. The executable generator is a relation validator, not a proof backend. The numerical Niederreiter set is a size anchor, not a deployment approval.

2 Notation and post-quantum conventions

The security parameter is λ . All adversaries and reductions are uniform or non-uniform QPT algorithms. A function $\text{negl}(\lambda)$ is negligible. Encodings are canonical, length-prefixed, and domain separated. Concatenation is written $\parallel$. Every hash invocation includes a fixed protocol label and version.

The issuing organization authenticates a real identity rid before the cryptographic issuance protocol. The session identifier sid is used only for replay and quota control and is never copied into a ticket. The common issuance information is

$$\text{ci} = (\text{version}, \text{epoch}, \text{domain}, \text{policy}, \text{expiryBucket}, \text{kid}_{\text{OA}}, \text{kid}_{\text{I}}).$$

Its fixed 32-byte ticket representation is

$$\text{ctx} = \text{SHAKE256}(\text{PQ-RBBC}/\text{CTX} \parallel \text{Encode}(\text{ci}), 256).$$

The federation configuration signature authenticates (ci, ctx) . The digest is computed once outside Π_{issue} and is a public input; the issuer cannot select a per-user value.

For any game G , $\text{Adv}_A^G(\lambda)$ denotes the usual absolute distance from the ideal winning probability. Oracle calls are polynomially bounded.

3 Required primitive interfaces

3.1 Hashing and CAP commitment

We use domain-separated hashes

$$H_{\text{ticket}}, \quad H_{\text{BUOV}}, \quad H_{\text{hold}}, \quad H_{\text{evidence}}.$$

They are modeled as quantum-resistant collision-resistant hashes. A random oracle or QROM instantiation may be used by a concrete component, but the core proof calls only the explicitly named property.

For the version-1.3 circuit audit, fix

$$\begin{aligned} H_{\text{ticket}}(X) &= \text{SHAKE256}(\text{PQ-RBBC}/\text{TICKET} \parallel X, 256), \\ H_{\text{hold}}(k) &= \text{SHAKE256}(\text{PQ-RBBC}/\text{HOLD} \parallel k, 256). \end{aligned}$$

The Blind-UOV hash H_{BUOV} and CAP commitment internals use the paper’s Anemoi-based costing profile; they are already included in its baseline NIZK count. This mixed choice retains standardized SHAKE/KMAC for new protocol logic while avoiding the multi-billion-constraint SHA3+AES baseline inside the Blind-UOV commitment proof.

There is no separate public commitment to the ticket payload. The accepting proof extracts M , Equation (I2) fixes its digest, and the Blind-UOV request binds that digest. Publishing an additional c_M would add a primitive and a circuit block without strengthening this chain.

3.2 Blind-UOV request normal form

The blind-signature interface is

$$(\text{bsk}, \text{bvk}) \leftarrow \text{BUOV.KeyGen}(1^\lambda),$$

with an interactive request/sign/unblind protocol and a public verification algorithm

$$\text{BUOV.Vfy}(\text{bvk}, m, \sigma).$$

The core needs more than ordinary blindness: the issuance proof must identify the exact hidden message represented by the request. The Blind-UOV authors' published construction admits the following paper-compatible normal form. Let $r \leftarrow \mathbb{F}_q^o$ be a fresh mask and let ρ_0 be commitment randomness. The user computes [1, 2]

$$c_0 \leftarrow \text{CAP.Commit}(r; \rho_0), \tag{B1}$$

$$y = r + H_{\text{BUOV}}(m, c_0) \in \mathbb{F}_q^o. \tag{B2}$$

The request object used by this core is $\beta = (c_0, y)$ and its blinding witness is $b = (r, \rho_0)$. The issuer applies the underlying UOV/hash-and-sign trapdoor to the masked target y only after accepting a proof that the request is well built. The user then runs the Blind-UOV finalization/CAP procedure to obtain the publicly verifiable final signature σ on m .

The exact byte-level CAP transcript is revision-dependent and is intentionally left behind the algorithms BUOV.Respond , BUOV.Finalize , and BUOV.Vfy . Equations (B1)–(B2) are the only Blind-UOV internals needed by the ticket-binding reduction. The authors' slides write CP.Commit ; this document writes CAP.Commit for that commitment subroutine inside the CAP system.

Definition 3.1 (Concrete Blind-UOV request relation). *For $\beta = (c_0, y)$ and $b = (r, \rho_0)$, define*

$$R_{\text{req}}^{\text{BUOV}}(\beta; m, b) = 1 \iff c_0 = \text{CAP.Commit}(r; \rho_0) \wedge y = r + H_{\text{BUOV}}(m, c_0).$$

A completed accepted protocol must return a Blind-UOV final signature that verifies on the same m .

We require correctness, blindness against an honest-protocol curious signer, and one-more unforgeability against malicious users in the QROM/security model of the selected Blind-UOV revision. Request binding is no longer an independent assumption: Lemma 7.2 derives it from CAP binding and collision resistance of H_{BUOV} .

Remark 3.2 (Why ordinary blind signing is insufficient). *If Π_{issue} proves a valid ciphertext but the proof is not tied to the exact blind request β , a malicious holder can prove statement M_0 and ask for a signature on M_1 . The $R_{\text{req}}^{\text{BUOV}}$ conjunct is the formal connection that prevents this substitution.*

3.3 Simulation-extractable NIZK argument of knowledge

The proof system $\Pi_{\text{issue}} = (\text{Setup}, \text{Prove}, \text{Vfy})$ is a post-quantum simulation-extractable NIZK argument of knowledge for the issuance relation in Section 4. Its setup produces a CRS; the security proof may use simulation and extraction trapdoors. In the real protocol no trapdoor is available to any party.

The proof may be several megabytes because it is transmitted only during infrequent issuance. This design choice is intentional. Knowledge soundness alone is enough for decoder safety in a simple stand-alone execution; simulation extractability is used to support concurrent sessions and the gated-CCA hybrid proof.

3.4 Compact threshold trace encryption

The threshold encryption interface is

$$\begin{aligned} (\text{tpk}, \text{tsk}_1, \dots, \text{tsk}_n) &\leftarrow \text{TEnc.KeyGen}(1^\lambda, n, \tau), \\ C &\leftarrow \text{TEnc.Enc}(\text{tpk}, \mu; \text{ad}, \rho), \\ d_i &\leftarrow \text{TPDec}(\text{tsk}_i, C), \\ \mu \text{ or } \perp &\leftarrow \text{TComb}(C, \{d_i\}_{i \in I}), \quad |I| \geq \tau. \end{aligned}$$

There is an NP relation

$$R_{\text{enc}}(\text{tpk}, C, \text{ad}; \mu, \rho) = 1 \iff C = \text{TEnc.Enc}(\text{tpk}, \mu; \text{ad}, \rho).$$

Only ciphertexts in the language L_{enc} are promised to be safe inputs to the distributed decoder.

We require the following post-quantum properties.

- (T1) **Correctness on valid inputs.** Combining any τ correct shares for a ciphertext in L_{enc} returns its plaintext.
- (T2) **τ -privacy under chosen plaintext.** A QPT adversary holding fewer than τ secret shares cannot distinguish encryptions of two equal-length plaintexts.
- (T3) **Unique opening.** It is infeasible to find $C, \mu, \rho, \mu', \rho'$ with $\mu \neq \mu'$ such that both $R_{\text{enc}}(\text{tpk}, C, \text{ad}; \mu, \rho) = 1$ and $R_{\text{enc}}(\text{tpk}, C, \text{ad}; \mu', \rho') = 1$.
- (T4) **Robust shares.** Invalid or inconsistent partial decryptions are detected by TComb or by a component proof/MAC before a plaintext is accepted.

Ordinary IND-CCA security of the bare encryption is not required and is not claimed.

3.5 Opening authorization

An independent post-quantum authorization mechanism verifies

$$\text{GateAuthVfy}(\text{gavk}, d_M, \text{caselD}, H_{\text{evidence}}(E), \text{auth}) = 1,$$

where $d_M = H_{\text{ticket}}(\text{Encode}(M))$. The authorization binds a ticket digest, case identifier, evidence digest, expiry, and purpose. Its unforgeability and policy governance are orthogonal to trace-ciphertext privacy, but the binding is required to prevent using an authorization for a different ticket.

4 Ticket format and issuance relation

The hidden ticket payload is

$$M = (\text{ctx}, \text{sn}, h, C), \quad h = H_{\text{hold}}(k_{\text{hold}}), \quad C = \text{Ntr.Enc}(\text{tpk}, \text{ctx}, h; \text{rid}, \text{sn}, e).$$

For the v1.3 profile all fields are fixed width and

$$\text{Encode}(M) = \text{ctx} \parallel \text{sn} \parallel h \parallel s \parallel c_\mu \parallel \gamma \in \{0, 1\}^{2944},$$

namely $32 + 16 + 32 + 208 + 48 + 32 = 368$ bytes. No variable-length field or user-selected policy string occurs in the payload. The serial number $\text{sn} \leftarrow \{0, 1\}^{128}$ and holder key $k_{\text{hold}} \leftarrow \{0, 1\}^{256}$ are fresh. The serial appears both outside and inside the trace ciphertext so that an opener can detect payload splicing. The holder secret itself never appears in the ticket. Define

$$m = H_{\text{ticket}}(\text{Encode}(M)).$$

The public issuance statement and private witness are

$$\begin{aligned} x &= (\text{pp}, \text{ctx}, \text{sid}, \text{rid}, \beta), \\ w &= (M, r, \rho_0, k_{\text{hold}}, e). \end{aligned}$$

The relation $R_{\text{issue}}(x; w) = 1$ iff all of the following hold:

$$\text{TicketShape}(M, \text{ctx}) = 1, \quad (\text{I1})$$

$$m = H_{\text{ticket}}(\text{Encode}(M)), \quad (\text{I2})$$

$$R_{\text{req}}^{\text{BUOV}}(\beta; m, (r, \rho_0)) = 1, \quad (\text{I3})$$

$$M.h = H_{\text{hold}}(k_{\text{hold}}), \quad (\text{I4})$$

$$R_{\text{enc}}(\text{tpk}, M.C, \text{Encode}(\text{ctx}, M.\text{sn}, M.h); \text{rid} \| M.\text{sn}, e) = 1. \quad (\text{I5})$$

The real identity and session identifier are public to the issuer during issuance but are not part of M . Equation (I5) binds the hidden trace plaintext to the identity authenticated in that exact session.

5 Construction

5.1 Setup

$\text{Core.Setup}(1^\lambda, n, \tau)$ performs:

1. generate Π_{issue} parameters and the domain-separated hash parameters;
2. generate issuer keys $(\text{bsk}, \text{bvk}) \leftarrow \text{BUOV.KeyGen}(1^\lambda)$;
3. generate independent opening keys $(\text{tpk}, \text{tsk}_1, \dots, \text{tsk}_n) \leftarrow \text{TEnc.KeyGen}(1^\lambda, n, \tau)$;
4. publish (ci, ctx) authenticated by a federation configuration signature.

No party is allowed to derive bsk and any tsk_i from the same seed or key hierarchy.

5.2 Blind relation-bound issuance

For authenticated identity rid and fresh issuer-side sid :

1. The holder samples $\text{sn}, k_{\text{hold}}, e$ and computes

$$h = H_{\text{hold}}(k_{\text{hold}}), \quad C = \text{Ntr.Enc}(\text{tpk}, \text{ctx}, h; \text{rid}, \text{sn}, e)$$

according to Equations (N1)–(N7), constructs M , and computes $m = H_{\text{ticket}}(\text{Encode}(M))$.

2. It samples r, ρ_0 , computes

$$c_0 = \text{CAP.Commit}(r; \rho_0), \quad y = r + H_{\text{BUOV}}(m, c_0),$$

sets $\beta = (c_0, y)$, and produces

$$\pi_{\text{issue}} \leftarrow \Pi_{\text{issue}}.\text{Prove}(\text{crs}, x, w).$$

3. The issuer checks authenticated identity, epoch, quota, replay status, and

$$\Pi_{\text{issue}}.\text{Vfy}(\text{crs}, x, \pi_{\text{issue}}) = 1.$$

Only then does it execute $z_{\text{BUOV}} \leftarrow \text{BUOV.Respond}(\text{bsk}, y)$.

4. The holder computes

$$\sigma \leftarrow \text{BUOV.Finalize}(m, c_0, y, z_{\text{BUOV}}; r, \rho_0)$$

and accepts only if $\text{BUOV.Vfy}(\text{bvk}, m, \sigma) = 1$.

The proof π_{issue} is not an extra proof placed beside the Blind-UOV well-built-request proof. It is the strengthened replacement for that proof: the original request equations are retained and the ticket, identity, holder, and trace-ciphertext conjuncts are appended to the proved relation. Thus issuance sends one large relation proof, not two independent proofs.

The spendable ticket is $T = (M, \sigma)$. Neither β , π_{issue} , rid , nor sid is transmitted with T .

5.3 Ticket verification

$\text{Core.VerifyTicket}(T)$ recomputes $m = H_{\text{ticket}}(\text{Encode}(M))$, checks canonical encoding and the public common-information policy, then returns $\text{BUOV.Vfy}(\text{bvk}, m, \sigma)$.

This check proves neither public ciphertext validity nor identity correctness by itself. Those facts follow from the certified-issuance theorem and the trust assumption that the issuer signs only after accepting Π_{issue} .

5.4 Signature-gated threshold opening

For an opening request

$$Q = (T, E, \text{caseID}, \text{auth}),$$

each honest opening authority runs the following gate before touching its decoder:

1. require $\text{Core.VerifyTicket}(T) = 1$;
2. require the common epoch and key identifiers to be acceptable;
3. compute $d_M = H_{\text{ticket}}(\text{Encode}(M))$ and require a valid, unused, purpose-limited authorization bound to $(d_M, \text{caseID}, H_{\text{evidence}}(E))$;
4. only after all checks pass, output $d_i = \text{TPDec}(\text{tsk}_i, M.C)$, bound to the same request digest.

The combiner accepts a reconstructed plaintext $\text{rid}' \parallel \text{sn}'$ only if robust share verification succeeds and $\text{sn}' = M.\text{sn}$. Otherwise it outputs $\perp$.

Critical API rule. The implementation MUST NOT export $\text{TPDec}(\text{tsk}_i, C)$ as a callable production endpoint. The only exposed operation is $\text{Core.OpenShare}(\text{tsk}_i, Q)$, which internally performs the gate. This is a security boundary, not an application convention.

6 Security games

6.1 Certified-language soundness

The challenger maintains a log $\mathcal{L}$ of accepted issuance sessions. In the analysis it uses the extractor to associate each completed session with (rid, M, w) . An adversary wins **CertSound** if it outputs a ticket $T = (M, \sigma)$ for which $\text{Core.VerifyTicket}(T) = 1$ and either:

- (a) M is not associated with any completed accepted issuance in $\mathcal{L}$; or
- (b) every associated extraction fails one of Equations (I1)–(I5).

Distinct messages are counted by their canonical ticket digest. This game captures both a new signature and a proof/request substitution.

6.2 Issuer unlinkability

An honest-protocol curious issuer chooses or observes two authenticated identities $\text{rid}_0, \text{rid}_1$ that receive the same ci . Two honest holders execute issuance concurrently. The issuer obtains its full protocol views, after which it receives the two resulting tickets in random order and guesses the order. It may retain all issuance state but receives no opening shares. The advantage over $1/2$ is issuer-unlinkability advantage.

This game allows the issuer to know which identity participated in each issuance; it hides which later ticket came from which session. It cannot hide timing, a uniquely sized policy, or a per-user expiration bucket, which is why common metadata is a stated system requirement.

6.3 Gated chosen-ciphertext trace privacy

The game **t-GCCA** is parameterized by threshold τ . The adversary receives all public parameters and fewer than τ opening secret shares. It may run malicious-user issuance sessions for identities it is authorized to use, receiving completed tickets, and may submit opening queries for any validly authorized ticket other than the exact challenge ticket digest.

For the challenge it selects equal-length identities $\text{rid}_0, \text{rid}_1$ and common ticket metadata. The challenger samples $b \leftarrow \{0, 1\}$, forms a valid ticket whose trace ciphertext encrypts $\text{rid}_b \parallel \text{sn}^*$, and returns $T^* = (M^*, \sigma^*)$. The adversary wins by guessing b .

The HNCC issuer is not corrupted in this game because it necessarily observes the authenticated identity at issuance. HNCC privacy is instead covered by issuer unlinkability. The game models an outside observer, an evidence processor, and fewer than τ opening authorities.

Remark 6.1 (Why the game is not ordinary CCA). *The opening oracle accepts only issuer-certified ticket payloads, not arbitrary ciphertexts. Calling the result IND-CCA for bare TEnc would be incorrect. The security statement is protocol-specific: certification plus the gate turns chosen-ciphertext access into chosen-certified-ciphertext access.*

6.4 Trace soundness and non-frameability

An adversary wins **TraceSound** if a valid ticket produced from a completed accepted issuance authenticated as rid opens to $\text{rid}' \neq \text{rid}$, or if it opens with an external serial number different from the encrypted serial. The issuer is honest in following the protocol. The holder may be malicious.

7 Security results

7.1 Correctness

Theorem 7.1 (End-to-end correctness). *For an honest issuance authenticated as rid , an honestly produced ticket $T = (M, \sigma)$ verifies. Any authorized set of at least τ honest opening authorities reconstructs $\text{rid} \parallel M.\text{sn}$, except with the sum of the component correctness errors.*

Proof. Honest computation satisfies every conjunct of R_{issue} , so Π_{issue} accepts. Blind-signature correctness gives a verifying σ on $H_{\text{ticket}}(\text{Encode}(M))$. Equation (I5) gives a valid encryption of $\text{rid} \parallel M.\text{sn}$. Threshold-encryption correctness on valid inputs reconstructs that plaintext from any τ correct shares. The serial cross-check therefore succeeds. $\square$

7.2 Blind-signing message binding

Lemma 7.2 (Blind-UOV request uniqueness). *Let $\beta = (c_0, y)$. Suppose $R_{\text{req}}^{\text{BUOV}}(\beta; m, (r, \rho_0)) = 1$ and $R_{\text{req}}^{\text{BUOV}}(\beta; m', (r', \rho'_0)) = 1$ for $m \neq m'$. Then one can derive either two different openings of the same CAP commitment or a collision in H_{BUOV} .*

Proof. Both witnesses open c_0 . If $r \neq r'$, they violate CAP binding. Otherwise $r = r'$, and the two equations for the same y imply

$$H_{\text{BUOV}}(m, c_0) = y - r = H_{\text{BUOV}}(m', c_0),$$

which is a collision because $m \neq m'$. Thus the request identifies at most one hidden message, except with the stated bad-event probability. $\square$

Lemma 7.3 (Request/proof binding). *Suppose an issuance proof for β is accepting and extractable. Except with CAP-binding or hash-collision probability, the message signed through β is exactly $H_{\text{ticket}}(\text{Encode}(M))$ for the payload M in the extracted witness. A different payload cannot verify under the same signed digest except by finding a collision in H_{ticket} .*

Proof. Extraction yields $(M, r, \rho_0, \dots)$ satisfying Equations (I2)–(I3). Equation (I2) fixes m to the canonical ticket digest. Lemma 7.2 makes the message represented by β equal to this m . If a different canonical payload $M' \neq M$ verifies under that signed value, then $H_{\text{ticket}}(\text{Encode}(M')) = H_{\text{ticket}}(\text{Encode}(M))$, a collision. $\square$

Theorem 7.4 (Certified-language soundness). *For any QPT adversary $\mathcal{A}$, there exist QPT reductions $\mathcal{B}_1, \dots, \mathcal{B}_5$ such that*

$$\begin{aligned} \text{Adv}_{\text{Core}, \mathcal{A}}^{\text{CertSound}} &\leq \text{Adv}_{\Pi_{\text{issue}}, \mathcal{B}_1}^{\text{SE-AoK}} + \text{Adv}_{\text{BUOV}, \mathcal{B}_2}^{1\text{mUF}} + \text{Adv}_{\text{CAP}, \mathcal{B}_3}^{\text{bind}} \\ &\quad + \text{Adv}_{H_{\text{BUOV}}, \mathcal{B}_4}^{\text{coll}} + \text{Adv}_{H_{\text{ticket}}, \mathcal{B}_5}^{\text{coll}}. \end{aligned}$$

Consequently, every verifying ticket is associated with an accepted issuance witness satisfying R_{issue} , except with negligible probability.

Proof. Consider a verifying output (M, σ) . If its canonical digest was not the message of a completed blind-signing session, then the output is a one-more forgery. Otherwise consider the corresponding accepting issuance proof. Simulation extraction returns a witness or violates the SE-AoK property. Lemma 7.3 associates the signed digest with the extracted payload, except under the remaining binding and collision bad events. The extracted witness then satisfies every issuance conjunct, including ciphertext validity and identity binding. $\square$

Corollary 7.5 (Certified ciphertext validity). *If $\text{Core.VerifyTicket}(M, \sigma) = 1$, then $M.C \in L_{\text{enc}}$ and it has an opening to the authenticated issuance identity and the external serial number, except with the bound in Theorem 7.4.*

7.3 Decoder safety

Theorem 7.6 (Signature-gated decoder safety). *Assume honest opening authorities expose only Core.OpenShare and execute its checks in the specified order. The probability that an honest authority evaluates TPDec on a ciphertext outside L_{enc} is at most*

$$\text{Adv}_{\text{Core}}^{\text{CertSound}} + \text{Adv}_{\text{GateAuth}}^{\text{forge}} + \text{Adv}_{H_{\text{ticket}}}^{\text{coll}}.$$

Proof. Reaching TPDec implies that the ticket signature and the opening authorization both verified for the same ticket digest. By Corollary 7.5, a verifying ticket contains a ciphertext in L_{enc} . A mismatch between the ticket checked by the gate and the ciphertext passed to the decoder either violates the atomic API invariant or requires a digest collision. A request not authorized for that digest is an authorization forgery. Taking a union bound gives the result. $\square$

Remark 7.7. *This theorem explains what is gained by moving the validity proof out of the online ticket. The large issuance proof is amortized, while the small issuer signature becomes a certificate that the ciphertext was once proved valid. The theorem fails immediately if an implementation offers a raw partial-decryption endpoint.*

7.4 Issuer unlinkability

Theorem 7.8 (Issuer unlinkability). *For an honest-protocol curious issuer, common visible metadata, and two honest holders, there exist QPT reductions such that*

$$\text{Adv}_{\text{Core}}^{\text{unlink}} \leq \text{Adv}_{\text{BUOV}}^{\text{blind}} + 2\text{Adv}_{\Pi_{\text{issue}}}^{\text{zk}} + \text{Adv}_{\text{TEnc}}^{\text{t-IND-CPA}} + \text{Adv}_{H_{\text{ticket}}}^{\text{coll}}.$$

Proof. First replace each real issuance proof with a simulated proof; zero knowledge bounds the change. The remaining request/response views are blind-signature views on two hidden messages. Blindness permits swapping those messages without the issuer detecting the swap. Threshold-encryption privacy prevents the issuer, which holds no opening threshold, from matching a final trace ciphertext to either known issuance identity. Canonical digest collisions are excluded. After these replacements, the issuer's view is independent of the final permutation, so its success probability is $1/2$. $\square$

Operational qualifier. The theorem is cryptographic, not traffic-analytic. Issuance batching, fixed message sizes, delayed activation, and common metadata are required to keep timing and formatting from becoming a link tag.

7.5 Trace correctness and non-frameability

Theorem 7.9 (Trace soundness). *Under certified-language soundness and unique opening of TEnc, a malicious holder completing issuance while authenticated as rid cannot produce a verifying ticket that opens to $\text{rid}' \neq \text{rid}$ or to a different serial number, except with negligible probability.*

Proof. By Theorem 7.4, a verifying ticket is associated with an extracted accepted issuance witness. Equation (I5) states that its ciphertext opens to $\text{rid} \parallel M.\text{sn}$. If threshold opening returns a different plaintext while all robust-share checks pass, either threshold correctness fails or the ciphertext has two different valid openings. The latter violates unique opening. The explicit serial comparison catches any external/internal serial mismatch. $\square$

Remark 7.10 (Issuer framing). *The theorem assumes an honest-protocol issuer because the issuer learns rid and controls the signing key. A malicious issuer could directly sign a payload encrypting an arbitrary victim identity. Removing that trust requires a threshold blind issuer, an auditable issuance log with privacy-preserving accountability, or retaining a publicly checked proof. It is not solved by the present compact-ticket core.*

7.6 Gated chosen-ciphertext trace privacy

Theorem 7.11 (τ -GCCA trace privacy). *Assume fewer than τ opening authorities are corrupted. For every QPT adversary $\mathcal{A}$ in the game of Section 6.3, there exist QPT reductions $\mathcal{B}_i$ such that*

$$\begin{aligned} \text{Adv}_{\text{Core}, \mathcal{A}}^{\text{t-GCCA}} &\leq \text{Adv}_{\text{TEnc}, \mathcal{B}_0}^{\text{t-IND-CPA}} + \text{Adv}_{\Pi_{\text{issue}}, \mathcal{B}_1}^{\text{SE-AoK}} + \text{Adv}_{\text{BUOV}, \mathcal{B}_2}^{\text{1mUF}} \\ &\quad + \text{Adv}_{\text{TEnc}, \mathcal{B}_3}^{\text{unique}} + \text{Adv}_{\text{CAP}, \mathcal{B}_4}^{\text{bind}} \\ &\quad + \text{Adv}_{H_{\text{BUOV}}, \mathcal{B}_5}^{\text{coll}} + \text{Adv}_{H_{\text{ticket}}, \mathcal{B}_6}^{\text{coll}}. \end{aligned}$$

Proof. We use a sequence of games.

Game 0. This is the real t-GCCA game.

Game 1: extract accepted issuance sessions. For every accepted adversarial issuance proof, use the SE-AoK extractor to record

$$(d_M, M, \text{rid}, e, k_{\text{hold}}, r, \rho_0, \dots)$$

in a private table $\mathcal{L}$. Abort if an accepting proof has no valid witness. The difference from Game 0 is bounded by SE-AoK security. Request binding and hash collision resistance ensure that the final signed digest is the digest recorded in $\mathcal{L}$.

Game 2: reject unlogged verifying tickets. If the adversary submits a verifying ticket whose digest is neither the challenge digest nor in $\mathcal{L}$, abort. Such a ticket supplies an additional valid message/signature pair not obtained from a completed issuance, which yields a one-more forgery. Proof/request substitution and digest ambiguity are charged to the corresponding terms in the theorem.

Game 3: simulate the opening oracle from the issuance table. For every non-challenge verifying ticket, look up the extracted plaintext in $\mathcal{L}$ and return it after the authorization checks. This has the same result as threshold opening by correctness and unique opening, but it uses no honest opening secret share and never invokes a decryption oracle. A ticket that reuses the challenge ciphertext under a newly certified digest requires an accepted proof containing a valid opening of that ciphertext. Extraction then reveals the challenge plaintext. If the extracted opening disagrees with the challenge plaintext, it violates unique opening; if it agrees, the reduction can immediately guess the threshold-encryption challenge bit. Thus challenge re-certification does not create an unanswerable decryption query.

Game 4: replace the challenge encryption. Embed the τ -IND-CPA challenge of TEnc as $M^*.C$. The reduction is given fewer than τ shares exactly as in the base game. It samples the remaining visible ticket fields identically. Because it generated the issuer key and knows the complete challenge message m^* , it locally executes both sides of the genuine Blind-UOV protocol and obtains an exactly distributed final signature; it does not need the unknown trace-encryption randomness. The challenge issuance transcript is not given to the adversary in this game, and π_{issue} influences only the accept/reject gate, not the distribution of the response after acceptance. Therefore the reduction may internally bypass that gate when manufacturing the challenge ticket without changing any value visible to the adversary. All opening queries are answered from $\mathcal{L}$ as in Game 3. Therefore changing the encrypted identity from rid_0 to rid_1 changes the adversary's view by at most the τ -IND-CPA advantage of TEnc.

In the final game the challenge bit is independent of the adversary's view. Summing the hybrid differences proves the theorem. $\square$

Why the reduction is useful. The large issuance proof is not merely a validity check. Its extractor gives the simulator the plaintext and randomness of every non-challenge ciphertext that can ever obtain a valid issuer signature. Therefore the simulator can answer gated opening queries without a base decryption oracle. This is the precise reason a CPA-secure compact threshold encryption can suffice inside the gated protocol even though it would not suffice under an unrestricted CCA API.

7.7 One-more ticket unforgeability

Theorem 7.12 (Ticket one-more unforgeability). *After q completed accepted blind-issuance sessions, producing more than q distinct verifying ticket digests has probability at most the one-more unforgeability advantage of BUOV plus the CAP-binding and hash-collision terms.*

Proof. Map every completed issuance to the unique signed digest using Lemma 7.3. More distinct verifying digests than completed sessions directly form a one-more forgery unless two payloads share a digest or Lemma 7.2 fails. $\square$

8 Candidate concrete instantiations

8.1 Augmented well-built-request compiler for Blind-UOV

The Blind-UOV framework already requires the user to prove that its masked request is well built. We strengthen that proved language instead of adding a second independent proof. Let

$$R_{\text{wb}}(c_0, y; m, r, \rho_0) = 1$$

denote Equations (B1)–(B2). The original well-built-request language is replaced by

$$\begin{aligned} R_{\text{aug}}(\text{pp}, \text{ctx}, \text{sid}, \text{rid}, c_0, y; \\ M, r, \rho_0, k_{\text{hold}}, e) = 1 \quad \Longleftrightarrow \quad R_{\text{issue}}(x; w) = 1, \end{aligned}$$

where $\beta = (c_0, y)$ and Equation (I3) expands to $R_{\text{wb}} = 1$. The issuer executes the ordinary Blind-UOV response only after verifying a post-quantum SE-NIZK argument for R_{aug} .

Compiler interface. The integration requires only one hook in a Blind-UOV implementation:

$$\text{VerifyWellBuilt}(c_0, y, \pi_{\text{wb}}) \longrightarrow \Pi_{\text{issue}}.\text{Vfy}(x, \pi_{\text{issue}}).$$

The response and finalization algorithms remain unmodified. In particular, the spendable signature is the genuine Blind-UOV final object. It need not be, and is not assumed to be, distributed as an ordinary direct UOV signature.

Concrete circuit boundary. The implementation SHOULD expose the following fixed interface to the proof compiler:

$$\begin{aligned} \text{pub}_{\text{aug}} &= (\text{ctx}, \text{sid}, \text{rid}, \text{tpk}, \text{bvk}, c_0, y), \\ \text{wit}_{\text{aug}} &= (\text{sn}, k_{\text{hold}}, e, r, \rho_0). \end{aligned}$$

The circuit derives C , h , M , and m internally rather than accepting duplicate witness copies. Its constraint system is the conjunction

$$\mathcal{C}_{\text{aug}} = \mathcal{C}_{\text{shape}} \wedge \mathcal{C}_{\text{ticketHash}} \wedge \mathcal{C}_{\text{BUOVmask}} \wedge \mathcal{C}_{\text{holder}} \wedge \mathcal{C}_{\text{trace}},$$

where

$$\begin{aligned} \mathcal{C}_{\text{shape}} &: \text{TicketShape}(M, \text{ctx}) = 1, \\ \mathcal{C}_{\text{ticketHash}} &: m = H_{\text{ticket}}(\text{Encode}(M)), \\ \mathcal{C}_{\text{BUOVmask}} &: c_0 = \text{CAP.Commit}(r; \rho_0) \wedge y = r + H_{\text{BUOV}}(m, c_0), \\ \mathcal{C}_{\text{holder}} &: h = H_{\text{hold}}(k_{\text{hold}}), \\ \mathcal{C}_{\text{trace}} &: R_{\text{Ntr}}(\text{tpk}, C, \text{ctx}, h; \text{rid}, \text{sn}, e) = 1. \end{aligned}$$

The proof statement includes sid even though it is absent from the ticket. Consequently, replaying the same proof in a new issuance session changes the statement and fails verification, while the final credential gains no session link tag.

Theorem 8.1 (Security inheritance of the augmented issuance proof). *Assume Π_{issue} is post-quantum zero knowledge and simulation-extractable for R_{aug} . Replacing the original well-built-request proof by π_{issue} preserves Blind-UOV one-more unforgeability and honest-protocol signer blindness, up to the SE-NIZK, CAP-binding, and hash-collision terms already shown in Theorems 7.4 and 7.8.*

Proof. For one-more unforgeability, a forger against the augmented scheme is also a forger against the original Blind-UOV scheme: the reduction verifies the augmented proof itself and forwards only accepted (c_0, y) requests to the original signer. If that interface requires the native well-built proof, the reduction extracts (m, r, ρ_0) from the augmented proof and generates the native proof; extraction failure is charged to the SE-AoK term. The augmented relation is a restriction of the original well-built language, so every forwarded request is valid there. The final signature format and verifier are unchanged.

For blindness, replace the augmented proof by a simulated proof. The remaining transcript is an ordinary Blind-UOV transcript, to which its blindness game applies. The public authenticated identity is session information but is absent from the final ticket; threshold-encryption privacy hides it inside $M.C$. Any attempt to make the same masked request represent a different ticket digest triggers Lemma 7.2, and any two payloads with one ticket digest trigger a collision in H_{ticket} . $\square$

Concretization achieved in versions 1.1–1.3. The request and witness are now fixed:

$$\begin{aligned} \beta &= (c_0, y), & b &= (r, \rho_0), \\ c_0 &= \text{CAP.Commit}(r; \rho_0), & y &= r + H_{\text{BUOV}}(H_{\text{ticket}}(\text{Encode}(M)), c_0). \end{aligned}$$

Consequently, the earlier abstract “request-binding” assumption has been eliminated. It is replaced by the standard, auditable assumptions that the CAP commitment is binding and the two domain-separated hashes are collision resistant. Version 1.2 removes c_M : extraction already supplies M , the request fixes its digest, and ticket-hash collision resistance prevents a second payload from verifying under that digest.

Frozen Blind-UOV costing profile. Version 1.3 fixes the 31 October 2025 revision and its *Blind-UOV-Is*, *NIST-I*, *Shorter*, *TCitH* row. The paper gives $\bar{n}_r = 256$ bits, $\bar{n}_x = 640$ bits, a 66.6 KB public key, and a 3,772-byte final signature. Its NIST-I/Shorter TCitH parameters are $\tau = 12$, ten 2^{11} -leaf trees plus two 2^{10} -leaf trees, $T_{\text{open}} = 111$, an explicit proof-of-work of 9 bits, and total proof-of-work 14.0 bits [1]. These values minimize the stored and online object; issuance is intentionally the expensive side of the trade-off.

Correct backend separation. Blind-UOV has two different proof objects. Its CAP commitment and final CAP proof are carried by the final blind signature. Before the signer responds, the user separately sends a NIZK π_1 proving that the masked request and CAP commitment are well formed. Our π_{issue} replaces and augments this π_1 ; it does *not* replace the final CAP proof. Consequently the 3,772-byte final signature remains unchanged, while the five ticket/trace blocks increase only the issuance NIZK circuit.

For costing, use the paper’s illustrative Shockwave/Anemoi row. Before our extension, Blind-UOV-Is/Shorter is reported as approximately 16 million R1CS constraints, a 2.8 MB proof, 32 s proving, and 2.5 s verification. The paper explicitly treats these as estimates and notes that the benchmarked implementations do not provide the required zero-knowledge mode. It also does not establish the simulation-extractability required by our concurrent GCCA proof. Therefore Shockwave/Anemoi is a *costing backend*, not yet the security-frozen Π_{issue} implementation [1].

The multivariate commitment $c = (\mu + F(r), G(r))$ and relation proofs of Bouillaguet et al. are a promising later optimization because they are algebraically aligned with UOV, but they

add a newer MQ/bilinear-system assumption and are not the baseline for the first auditable construction [3].

8.2 Threshold Niederreiter for trace encryption

Code and plaintext profile. Let $\mathcal{G}$ be a binary Goppa code of length n , dimension k , and designed error weight t , with public systematic parity-check matrix $H_{\text{pub}} = (I_{n-k} \mid T)$. The selected decoder MUST guarantee $d_{\min}(\mathcal{G}) \geq 2t + 1$. Fix $\text{rid} \in \{0, 1\}^{256}$ and $\text{sn} \in \{0, 1\}^{128}$, and define the fixed 48-byte plaintext

$$\mu = \text{Encode}_{48}(\text{rid} \parallel \text{sn}).$$

The online ciphertext uses a weight- t error $e \in \mathbb{F}_2^n$ as its only encryption witness. Thus the abstract randomness ρ_{tr} is instantiated by e .

Hash DEM. Sample e uniformly from the weight- t vectors and compute

$$s = H_{\text{pub}} e^\top \in \mathbb{F}_2^{n-k}, \quad \text{wt}(e) = t, \quad (\text{N1})$$

$$\text{ad} = \text{Encode}(\text{ctx}, \text{sn}, h), \quad (\text{N2})$$

$$Z = \text{SHAKE256}(\text{PQ-RBBC}/\text{KDF} \parallel e \parallel s \parallel \text{ctx}, 640), \quad (\text{N3})$$

$$(K_{\text{mac}}, P) = Z, \quad |K_{\text{mac}}| = 32, \quad |P| = 48, \quad (\text{N4})$$

$$c_\mu = \mu \oplus P, \quad (\text{N5})$$

$$\gamma = \text{KMAC256}_{K_{\text{mac}}}(s \parallel c_\mu \parallel \text{ad}, 256; \text{custom} = \text{PQ-RBBC}/\text{TAG}), \quad (\text{N6})$$

$$C = (s, c_\mu, \gamma). \quad (\text{N7})$$

The output length in (N3) is in bits. SHAKE256 and KMAC256 are the standardized extendable-output and MAC functions of FIPS 202 and SP 800-185, respectively [8, 9]. The composition HDEM is nevertheless a new protocol profile, not a standardized AEAD and not yet a production claim. A standardized one-time AEAD may replace it later without changing the abstract R_{enc} interface.

Definition 8.2 (Concrete trace relation). *For $C = (s, c_\mu, \gamma)$, define*

$$R_{\text{Ntr}}(\text{tpk}, C, \text{ctx}, h; \text{rid}, \text{sn}, e) = 1$$

iff tpk selects H_{pub} and the fixed profile above, and Equations (N1)–(N7) all hold with $\mu = \text{Encode}_{48}(\text{rid} \parallel \text{sn})$. In this version,

$$\begin{aligned} R_{\text{enc}}(\text{tpk}, C, \text{Encode}(\text{ctx}, \text{sn}, h); \text{rid} \parallel \text{sn}, \rho_{\text{tr}}) &= 1 \\ \iff R_{\text{Ntr}}(\text{tpk}, C, \text{ctx}, h; \text{rid}, \text{sn}, e) &= 1, \quad \rho_{\text{tr}} = e. \end{aligned}$$

Hence $\mathcal{C}_{\text{trace}}$ is now the exact conjunction (N1)–(N7), not an unspecified “encrypt correctly” box.

Lemma 8.3 (Unique syndrome witness). *For a code with $d_{\min} \geq 2t + 1$, a syndrome s has at most one witness e of weight at most t satisfying $s = H_{\text{pub}} e^\top$.*

Proof. If e and e' were two such witnesses, then $d = e + e'$ would satisfy $H_{\text{pub}} d^\top = 0$ and $\text{wt}(d) \leq 2t$. If $e \neq e'$, d is a nonzero codeword below the minimum distance, a contradiction. Therefore $e = e'$. This is the standard unique-decoding argument used by Classic McEliece [6]. $\square$

Corollary 8.4 (Unique trace opening). *The concrete relation R_{Ntr} has information-theoretically unique plaintext openings.*

Proof. For a fixed C , Lemma 8.3 fixes e from s . Equations (N3)–(N4) then deterministically fix P and K_{mac} , and (N5) fixes the sole $\mu = c_\mu \oplus P$. The fixed-width encoding is injective, so it fixes (rid, sn) . Equation (N6) additionally rejects any inconsistent associated data. No computational assumption is needed for uniqueness. $\square$

Proposition 8.5 (Concrete trace correctness). *If threshold decoding reconstructs the unique error e for s , the combiner recomputes (N3)–(N6), returns $\mu = c_\mu \oplus P$, and accepts only if γ verifies and the decoded serial equals the visible sn . Consequently every honestly formed concrete trace ciphertext opens to exactly $\text{rid} \parallel \text{sn}$.*

Privacy assumption and proof boundary. Let t-OW-SD be the game in which an adversary receives H_{pub} , fewer than τ threshold shares, and $s = H_{\text{pub}}e^\top$ for uniform weight- t e , and must output e . We assume this advantage is negligible for the frozen distributed decoder.

Theorem 8.6 (Conditional concrete trace privacy). *Model SHAKE256 as a quantum random oracle, assume KMAC256 is a quantum-secure PRF, never reuse e , and let $\epsilon_{\text{O2H}}(q_H)$ denote the loss of the applicable one-way-to-hiding lemma for q_H quantum hash queries. Then for every fewer-than-threshold QPT distinguisher $\mathcal{A}$ there are QPT reductions $\mathcal{B}_1, \mathcal{B}_2$ such that*

$$\text{Adv}_{\text{Ntr}, \mathcal{A}}^{\text{t-IND-CPA}} \leq \text{Adv}_{\mathcal{G}, \mathcal{B}_1}^{\text{t-OW-SD}} + \epsilon_{\text{O2H}}(q_H) + \text{Adv}_{\text{KMAC256}, \mathcal{B}_2}^{\text{qPRF}}.$$

Proof. Starting from the real game, replace the 640-bit SHAKE256 value at the hidden input (e, s, ctx) by a uniform string. A distinguisher for this step either yields a threshold syndrome-decoding solver or is bounded by the stated one-way-to-hiding loss. Now P is an independent uniform 48-byte pad and K_{mac} is an independent uniform key. Replace KMAC256 under that key by a random function, paying the qPRF term. The value $c_\mu = \mu_b \oplus P$ and the tag are then independent of b , so the final advantage is zero. Summing the hybrid differences gives the bound. $\square$

A new accepted tuple (s, c'_μ, γ') with the same hidden derived key and different authenticated input gives a KMAC forgery, up to the same syndrome-decoding/QROM bad event. Thus the xor pad gives one-time privacy and KMAC256 gives one-time ciphertext integrity. Substituting Theorem 8.6 for the abstract τ -IND-CPA term in Theorem 7.11 yields protocol-level τ -GCCA, plus the SE-AoK, Blind-UOV, CAP, unique-opening, authorization, and hash terms already displayed there. This does *not* prove standard IND-CCA security for a bare ciphertext endpoint.

The threshold-Niederreiter work solves malformed-ciphertext attacks by attaching a straight-line-extractable validity proof to every public ciphertext; its authors report about 30.5 KB for a three-party ciphertext, with the proof dominating the syndrome [4, 5]. Here the same validity language $\{s : \exists e, \text{wt}(e) = t, s = He^\top\}$ is certified once inside π_{issue} . The online ticket retains only $C = (s, c_\mu, \gamma)$, and the raw distributed decoder remains unreachable.

Exact online byte counts. Because $|c_\mu| = 48$ and $|\gamma| = 32$, the trace ciphertext is exactly the syndrome length plus 80 bytes. Using the published Classic McEliece implementation sizes gives the following anchors [7]:

Parameter anchor	Public key	Syndrome	This C_{trace}
mceliece348864	261120 B	96 B	176 B
mceliece460896	524160 B	156 B	236 B
mceliece6688128	1044992 B	208 B	288 B
mceliece6960119	1047319 B	194 B	274 B
mceliece8192128	1357824 B	208 B	288 B

The aggressive working anchor is mceliece6688128: approximately 1.045 MB of cached public key buys a 288-byte trace ciphertext. The frozen Blind-UOV-Is profile and exact 4,140-byte complete ticket are derived in Section 8.3; these are still research targets rather than implementation benchmarks.

Remark 8.7 (Mandatory 2026 cryptanalytic hold). *The anchor above is not frozen. A 10 August 2026 preprint gives a classical quasipolynomial-time distinguisher applying to all Classic McEliece NIST parameter sets and extends it to a heuristic quasipolynomial-time decryption attack; its concrete improvements are described as not yet practical [10]. That is not a demonstrated practical break, but it materially changes confidence. Deployment parameter selection is blocked until independent review quantifies this attack for the threshold construction.*

8.3 Reproducible five-block circuit audit

Counting model. The paper’s Anemoi profile uses an R1CS over $\mathbb{F}_{2^{129}}$ for NIST level I. Addition, xor, multiplication by public constants, fixed-width concatenation, and the public binary syndrome map are linear. Each Boolean AND or field multiplication contributes one nonlinear R1CS constraint. Following the Blind-UOV paper, one Keccak-f permutation costs $24 \cdot 1600 = 38,400$ constraints [1]. The audit counts only constraints added beyond the paper’s approximately 16-million-constraint Blind-UOV-Is/Shorter π_1 baseline.

Exact-weight gadget. The witness e has 6,688 bits. Pad it with 1,504 public zeros to 8,192 inputs and sum them using a fixed balanced binary adder tree. Combining two j -bit counters uses one half-adder and $j - 1$ full adders. Since xor is linear, the nonlinear cost is $1 + 2(j - 1) = 2j - 1$. The exact cost is therefore

$$\sum_{j=1}^{13} \frac{8192}{2^j} (2j - 1) = 24,547$$

AND constraints. Equating the final 14-bit count to 128 is linear. This gadget avoids the invalid shortcut $\sum e_i = 128$ in characteristic two, which would check only parity. Separately, e , sn , and k_{hold} require $6688 + 128 + 256 = 7,072$ bitness constraints.

Exact Keccak accounting. SHAKE256 has a 136-byte rate. With pad_{10^*1} , an L -byte message uses $\lfloor L/136 \rfloor + 1$ absorb permutations. The fixed encodings give:

Invocation	Absorbed bytes	Keccak-f calls	Nonlinear constraints
$H_{\text{ticket}}(M)$	$14 + 368 = 382$	3	115,200
$H_{\text{hold}}(k)$	$12 + 32 = 44$	1	38,400
Trace SHAKE KDF	$11 + 836 + 208 + 32 = 1087$	8	307,200
KMAC256	$136 + 136 + 336 + 3 = 611$	5	192,000
Total		17	652,800

For KMAC, the first 136-byte block is the cSHAKE function/customization prefix, the second is $\text{bytepad}(\text{encode_string}(K), 136)$, the authenticated data is $208 + 48 + 80 = 336$ bytes, and $\text{right_encode}(256)$ is three bytes. All squeezing fits in the first rate block and adds no permutation.

Five-block result. The incremental nonlinear counts are:

Block	Constraints	Contents
$\mathcal{C}_{\text{shape}}$	128	serial bitness; widths/wiring linear
$\mathcal{C}_{\text{ticketHash}}$	115,200	three SHAKE permutations
$\mathcal{C}_{\text{BUOVmask}}$	0	already in the paper baseline
$\mathcal{C}_{\text{holder}}$	38,656	256 bitness + one SHAKE permutation
$\mathcal{C}_{\text{trace}}$	530,435	e bitness, weight, KDF, and KMAC
Added relation	684,419	exact for the specified gadgets

The syndrome equation adds 1,664 linear output equalities but no nonlinear gate. The xor pad, associated-data construction, external/internal serial equality, and final constant comparisons are also linear.

Paper-anchored issuance estimate. Because the paper rounds its baseline to 16 million constraints, the combined number is itself approximate:

$$N_{\pi_{\text{issue}}} \approx 16,000,000 + 684,419 = 16,684,419,$$

an increase of 4.2776%. Applying the paper’s square-root proof-size scaling and linear prover scaling gives approximately 2.86 MB and 33.4 s. A conservative linear scaling of its reported verifier figure gives 2.61 s. These are extrapolations from Table 5 of the Blind-UOV paper, not local benchmarks. The paper does not provide enough information to derive peak prover memory, so memory remains a measurement obligation.

Proposition 8.8 (Monolithic-circuit feasibility decision). *Under the costing model above, adding the entire ticket and trace relation to the existing well-built-request NIZK changes its constraint count by less than 5%. There is therefore no communication-based reason to split R_{aug} into linked proofs at this stage. A split is justified only if the selected simulation-extractable backend, implementation memory, or prover parallelism makes the monolithic circuit impractical.*

Exact online object under the frozen profile. For mceliece6688128, $|C| = 288$ bytes and

$$|M| = 32 + 16 + 32 + 288 = 368 \text{ bytes.}$$

Combining this with the published 3,772-byte Blind-UOV-Is/Shorter signature gives

$$|T| = |M| + |\sigma| = 4,140 \text{ bytes}$$

before transport framing. The cached public material is approximately 1,044,992 bytes for the trace key plus the paper’s rounded 66.6 KB Blind-UOV verification key. Relative to a 70 KB online credential, the frozen target is about 17 times smaller. The approximately 2.86 MB issuance NIZK is offline and is never forwarded to a satellite verifier.

Remark 8.9 (Audit reproducibility). *The companion script `pq_rbbc_circuit_audit.py` recomputes every integer above and emits the checked `pq_rbbc_circuit_audit_v1_3.json`. Exact gadget counts and paper-derived estimates are kept in separate JSON fields so a rounded source number cannot be mistaken for a benchmark.*

8.4 Executable reference relation and rejection tests

Version 1.4 implements the incremental relation in `pq_rbbc_reference.py`. The witness builder derives h , s , c_μ , γ , the canonical 368-byte payload, and its ticket digest from $(\text{rid}, \text{ctx}, \text{sn}, k_{\text{hold}}, e)$. The circuit builder then emits a stream of characteristic-two events of four kinds: linear wire definitions, linear equality assertions, Boolean multiplication gates, and bitness constraints

$x(x + 1) = 0$. A backend can consume the same stream without retaining the whole circuit in memory. This version executes the stream and counts it, but does not yet flatten it into a proof system’s matrix format.

The symbolic Keccak-f implementation performs all 24 rounds. Its SHAKE256 output is checked against Python’s independent FIPS 202 implementation, and the SP 800-185 KMAC256 output is checked against the OpenSSL 3 KMAC provider [8, 9]. For the frozen honest vector the generated circuit reports:

Measured generator quantity	Value
Allocated wire identifiers	2,976,528
Linear definitions	2,281,323
Linear equality assertions	3,214
Nonlinear constraints	684,419
Keccak-f permutations	17
Failed assertions	0

Thus the executable result exactly reproduces every nonlinear block total in the preceding audit. The linear-definition figure is diagnostic IR size, not an additional nonlinear count: the costing model requires a backend to inline affine expressions. A backend that instead materializes each affine wire as an R1CS row must report that larger total explicitly. On the present reference environment, one honest streaming generation took approximately 5.85 seconds. This is Python relation generation time, not NIZK proving time, and no prover-memory claim follows from it.

The regression suite accepts the honest witness and sends eight independent mutations through the complete circuit: wrong error weight, changed syndrome, changed masked identity, changed holder hash, changed KMAC tag, changed ticket digest, changed blind request, and changed common context. All eight are rejected. The suite contains eight test methods and completed in approximately 53.5 seconds in the same environment. The checked output is recorded in `pq_rbbc_r1cs_manifest_v1_4.json`.

Exact implementation boundary. The $H = (I \mid T)$ matrix in the reference vector is generated deterministically from a public seed and is only a rank-controlled test fixture; it is not a certified binary Goppa key. Similarly, `TestBlindUOVAdapter` checks message/request binding with a deterministic SHAKE construction, but is not a blind signature. It marks the native Blind-UOV π_1 /CAP boundary whose nonlinear cost remains in the paper’s baseline. Consequently this implementation establishes that the incremental relation is internally consistent and executable; it does not establish that the combined native Blind-UOV proof is zero knowledge or simulation extractable.

8.5 Local mixing as a future trace-encryption candidate

The core does not depend on elliptic curves, lattices, or codes. A future local-mixing encryption may replace threshold Niederreiter only if it supplies all four properties (T1)–(T4), an efficiently provable NP encryption relation, and a practical τ -party opening protocol. A small standalone ciphertext claim is insufficient by itself: threshold opening, unique opening, and extractable issuance binding are the missing system properties.

9 What has and has not been proved

Status	Result or obligation
Proved at the abstract level	Correctness; proof/request/message binding; certified ciphertext validity; signature-gated decoder safety; issuer unlinkability against an honest-protocol curious issuer; trace soundness; one-more ticket unforgeability; and τ -GCCA trace privacy.
Assumed component property	Post-quantum SE-NIZK AoK, CAP binding, Blind-UOV blindness and one-more unforgeability, threshold syndrome-decoding privacy with robust shares, QROM one-way-to-hiding, KMAC integrity, and authorization unforgeability.
Concretely instantiated	Blind-UOV-Is/Shorter costing profile; executable five-block augmented relation; exact Niederreiter-SHAKE-KMAC trace relation; information-theoretic unique opening; independently checked SHAKE/KMAC; full-circuit positive and negative tests; 684,419 added nonlinear constraints; and a 4,140-byte online-ticket target.
Not yet instantiated	A post-quantum zero-knowledge and simulation-extractable implementation of Π_{issue} ; native Blind-UOV circuit integration; backend matrix lowering; import of a real Goppa key; peak prover memory; robust threshold-share transcript; implementation benchmarks; and a reviewed response to the August 2026 McEliece cryptanalysis.
Explicitly not claimed	Standard IND-CCA security of the bare ciphertext, security against a malicious issuer, metadata/timing anonymity without system controls, UC composability, side-channel resistance, satellite AKE security, or FLEO/handover security.

10 Implementation invariants

The following invariants are part of the cryptographic design and should be machine-tested or enforced by type/API boundaries.

- INV-1 There is no production function that accepts a bare trace ciphertext and returns a partial decryption.
- INV-2 The bytes hashed for the issuer signature are exactly the canonical bytes parsed by ticket verification and opening.
- INV-3 The issuer signs only after successful proof verification and records a completed session only after returning its blind response.
- INV-4 The blind request proven in Π_{issue} is byte-for-byte the request passed to the blind-signing implementation.
- INV-5 The opening authorization includes the ticket digest and evidence digest, and replay is statefully rejected.
- INV-6 The reconstructed serial number must equal the clear serial number.

- INV-7 Issuer keys and opening keys have independent generation ceremonies, storage, rotation, and compromise domains.
- INV-8 Epoch metadata is federation-fixed; the issuer cannot personalize it.
- INV-9 Every component and parameter set has an explicit post-quantum security claim at the chosen target category.
- INV-10 The parser accepts exactly the frozen 368-byte payload layout and rejects variable-length or alternate encodings before signature verification.
- INV-11 The issuance circuit enforces bitness of all 6,688 error coordinates and proves exact weight 128 with the specified counter gadget; a parity-only check is forbidden.
- INV-12 The Shockwave/Anemoi figures may be used for costing only until a backend meeting the stated zero-knowledge and simulation-extractability model is proved or substituted.
- INV-13 Production configuration rejects the deterministic parity-check matrix and test-only blind-signature adapter; neither fixture may cross the backend boundary.
- INV-14 Backend lowering must preserve every multiplication and equality event, eliminate affine wires soundly, and accept/reject the published reference and mutation vectors identically to the executable IR.

11 Next proof step

The relation, unique opening, fixed layout, circuit audit, executable witness builder, and rejection tests are now complete. The next step is native composition and proof-backend qualification:

Lower the streaming relation to the selected backend, replace the test Blind-UOV adapter with the native π_1 /CAP constraints, and prove that the composed Π_{issue} is post-quantum zero knowledge and simulation-extractable in the model required by the GCCA reduction. Then import a canonical Goppa public key and measure prover time and peak memory.

The current count favors one monolithic proof; splitting is no longer the default. The Niederreiter deployment parameter MUST remain provisional under Remark 8.7. After backend qualification, the remaining cryptographic component is the robust threshold-share transcript. Satellite AKE integration remains out of scope until those two obligations close.

References

- [1] Charles Bouillaguet, Thibauld Feneuil, Jules Maire, Matthieu Rivain, Julia Sauvage, and Damien Vergnaud, *Blinding Post-Quantum Hash-and-Sign Signatures*, IACR Cryptology ePrint Archive, Report 2025/895, 2025. <https://eprint.iacr.org/2025/895>
- [2] Thibauld Feneuil, *Blinding Post-Quantum Hash-and-Sign Signatures*, IEEE Symposium on Security and Privacy presentation slides, 2026. https://www.thibauld-feneuil.fr/docs/talks/2026-05-19_SP26_blind-sigs.pdf
- [3] Charles Bouillaguet, Thibauld Feneuil, Jules Maire, Matthieu Rivain, Julia Sauvage, and Damien Vergnaud, *Multivariate Commitments and Signatures with Efficient Protocols*, IACR Cryptology ePrint Archive, Report 2025/2035, 2025. <https://eprint.iacr.org/2025/2035>

- [4] Pascal Giorgi, Fabien Laguillaumie, Lucas Ottow, and Damien Vergnaud, *Threshold Niederreiter: Chosen-Ciphertext Security and Improved Distributed Decoding*, IACR Cryptology ePrint Archive, Report 2025/757, 2025. <https://eprint.iacr.org/2025/757>
- [5] Lucas Ottow, *Threshold Niederreiter: Chosen-Ciphertext Security and Improved Distributed Decoding*, presentation slides, 2025. https://indico.math.cnrs.fr/event/11948/attachments/5546/9681/Lucas_Ottow.pdf
- [6] Daniel J. Bernstein et al., *Classic McEliece: conservative code-based cryptography*, specification. <https://classic.mceliece.org/mceliece-spec-20221023.pdf>
- [7] Classic McEliece team, *Implementations and parameter sizes*. <https://classic.mceliece.org/impl.html>
- [8] National Institute of Standards and Technology, *FIPS 202: SHA-3 Standard*, 2015. <https://csrc.nist.gov/pubs/fips/202/final>
- [9] National Institute of Standards and Technology, *SP 800-185: SHA-3 Derived Functions*, 2016. <https://csrc.nist.gov/pubs/sp/800/185/final>
- [10] Ashrujit Ghoshal, Yuval Ishai, Aayush Jain, and Nuozhou Sun, *Quasipolynomial Cryptanalysis of the McEliece Cryptosystem (or: PIR Meets McEliece)*, IACR Cryptology ePrint Archive, Report 2026/1630, 2026. <https://eprint.iacr.org/2026/1630>
- [11] Olivier Blazy, Georg Fuchsbauer, David Pointcheval, and Damien Vergnaud, *Signatures on Randomizable Ciphertexts*, Public Key Cryptography—PKC 2011. https://www.blazy.eu/Papers/2011_pkc.pdf